

```

from typing import Optional, Tuple
import torch
from torch import nn
from torch.nn import functional as F
import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm
import os
import json
#import imageio
import cv2
import os
os.environ['PYTORCH_CUDA_ALLOC_CONF'] = 'expandable_segments:True'
from PIL import Image

from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly



```

def load_colmap_data():
    r"""
    After using colmap2nerf.py to convert the colmap intrinsics and ex
    read in the transform_colmap.json file

    Expected Returns:
        An array of resized imgs, normalized to [0, 1]
        An array of poses, essentially the transform matrix
        Camera parameters: H, W, focal length

    NOTES:
        We recommend you resize the original images from 800x800 to lowe
        i.e. 200x200 so it's easier for training. Change camera paramete
    """
    ##### YOUR CODE START #####

```

```

json_path='transforms_colmap.json'
image_dir='./data/data/images'

json_path='/content/drive/MyDrive/ESE 650_HW3_P3/transforms_colmap'
image_dir='/content/drive/MyDrive/ESE 650_HW3_P3/data/data/images'
resize_dim=(200, 200)

# Load JSON file
with open(json_path, 'r') as f:
    data = json.load(f)

# Initialize lists to store processed data
imgs = []
poses = []

# Process each frame in the JSON file
for frame in data['frames']:
    # Load and resize image
    img_path = os.path.join(image_dir, frame['file_path'][0])
    img = Image.open(img_path).resize(resize_dim)
    imgs.append(np.array(img) / 255.0) # Normalize to [0, 1]

    # Process pose
    pose = np.array(frame['transform_matrix'])
    #print (pose)
    poses.append(pose)

# Assuming all images have the same size and focal length
H, W = resize_dim
focal = frame["camera_angle_x"] * W / (2 * np.tan(frame['camera_ar

# Convert lists to numpy arrays
imgs = np.array(imgs)
poses = np.array(poses)

# Camera parameters: Height, Width, Focal length
camera_params = (H, W, focal)

return imgs, poses, camera_params
##### YOUR CODE END #####

```

```

def get_rays(H, W, focal, c2w,):
    """
    Compute rays passing through each pixel in the world frame.

    Parameters:
    - H: Height of the image.
    - W: Width of the image.
    - focal: Camera intrinsic matrix of shape (3, 3).
    - c2w: Camera-to-world transformation matrix of shape (4, 4).

    Returns:
    - ray_origins: Array of shape (H, W, 3) denoting the origins of each ray.
    - ray_directions: Array of shape (H, W, 3) denoting the direction of each ray.
    """
    # Generate mesh grid for pixel coordinates
    i, j = torch.meshgrid(torch.linspace(0, W-1, W, device=device), torch.linspace(0, H-1, H, device=device))
    i = i.t().flatten()
    j = j.t().flatten()

    # Detect if a GPU is available and choose the device accordingly
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    #device = 'cpu' #Due to OOM
    #print(device)

    # Ensure c2w is a torch.Tensor and move it to the chosen device
    if not isinstance(c2w, torch.Tensor):
        c2w = torch.from_numpy(c2w).float() # Convert from numpy to torch
    c2w = c2w.to(device)

    # Generate a grid of (i, j) coordinates
    i, j = torch.meshgrid(torch.linspace(0, W-1, W, device=device), torch.linspace(0, H-1, H, device=device))
    i = i.t().flatten()
    j = j.t().flatten()

    # Normalize pixel coordinates (assuming the image center as origin)
    dirs = torch.stack([(i - W * 0.5) / focal, -(j - H * 0.5) / focal, 1.0])

    # Rotate ray directions from camera frame to the world frame
    rays_d = torch.sum(dirs[None, :, :3] * c2w[:3, :3], axis=-1)

    # The origin of all rays is the camera position in the world frame
    rays_o = c2w[:3, -1].expand(rays_d.shape)

```

```

# Reshape rays_o and rays_d to [H, W, 3]
rays_o = rays_o.view(H, W, 3)
rays_d = rays_d.view(H, W, 3)

return rays_o, rays_d

def sample_points_from_rays(ray_origins, ray_directions, snear, sfar,
    """
    Sample 3D points along rays within the specified near and far bound

    Parameters:
    - ray_origins: Array of shape (H, W, 3) denoting the origins of each ray
    - ray_directions: Array of shape (H, W, 3) denoting the direction of each ray
    - snear: Scalar or array defining the near clipping distance for each ray
    - sfar: Scalar or array defining the far clipping distance for each ray
    - Nsample: Number of points to sample along each ray.

    Returns:
    - sampled_points: Array of shape (H, W, Nsample, 3) with sampled 3D points
    - depth_values: Array of shape (H, W, Nsample) with depth values corresponding to each sample
    """

    # Make sure snear and sfar are tensors
    H, W, _ = ray_origins.shape
    device = ray_origins.device

    # Compute the depth values for each sample
    depth_values = torch.linspace(snear, sfar, Nsample, device=device)

    depth_values = depth_values.expand(H, W, Nsample) # Make sure depth values are broadcastable

    # Compute the 3D positions of each sample point along the rays
    sampled_points = ray_origins[..., None, :] + depth_values[..., :, None] * ray_directions

    return sampled_points, depth_values

def positional_encoding(x, max_freq_log2=10, include_input=True):
    """Apply positional encoding to the input (Section 5.1 of original paper)"""

```

Apply positional encoding to the input. (Section 5.1 of original paper)
We use positional encoding to map continuous input coordinates into higher dimensional space to enable our MLP to more easily approximate higher frequency function.

Expected Returns:

pos_out: positional encoding of the input tensor.
(H*W*num_samples, (include_input + 2*freq) * 3)

"""

```
frequencies = 2 ** torch.linspace(0, max_freq_log2, steps=max_freq_log2)
# Create a list of frequencies, (sin(2^k * x), cos(2^k * x)) for k in range(max_freq_log2)
encodings = [torch.sin(x * freq) for freq in frequencies] + [torch.cos(x * freq) for freq in frequencies]
# Stack all encodings along the last dimension
encoded = torch.cat(encodings, dim=-1)
```

```
if include_input:
```

```
    # Concatenate the original input with the encoded features
    pos_out = torch.cat([x, encoded], dim=-1)
```

```
else:
```

```
    pos_out = encoded
```

```
return pos_out
```

```
def volume_rendering(
    radiance_field: torch.Tensor,
    ray_origins: torch.Tensor,
    depth_values: torch.Tensor
```

```
) -> Tuple[torch.Tensor]:
```

"""

Differentiably renders a radiance field, given the origin of each ray and the sampled depth values along them.

Args:

radiance_field: Tensor containing RGB color and volume density at each point, shape (H, W, num_samples, 4).

ray_origins: Origin of each ray, shape (H, W, 3).

depth_values: Sampled depth values along each ray, shape (H, W, num_samples).

Returns:

rgb_map: Rendered RGB image, shape (H, W, 3).

```

"""
# Extract sigma (density) and color from the radiance field
sigma = torch.relu(radiance_field[..., 3]) # Extract volume density
rgb = torch.sigmoid(radiance_field[..., :3]) # Extract RGB colors

# Compute depth intervals
dists = torch.cat([depth_values[..., 1:] - depth_values[..., :-1],
alpha = 1.0 - torch.exp(-sigma * dists)
weights = alpha * torch.cumprod(torch.cat([torch.ones_like(alpha[.

rgb_map = torch.sum(weights[..., None] * rgb, dim=-2)

return rgb_map

class TinyNeRF(torch.nn.Module):
    def __init__(self, pos_dim, fc_dim=128):
        """Initialize a tiny nerf network, which composed of linear lay
        ReLU activation. More specifically: linear - relu - linear - rel
        - relu -linear. The module is intentionally made small so that w
        achieve reasonable training time

        Args:
            pos_dim: dimension of the positional encoding output
            fc_dim: dimension of the fully connected layer
        """
        super().__init__()

        self.nerf = nn.Sequential(
            nn.Linear(pos_dim, fc_dim),
            nn.ReLU(),
            nn.Linear(fc_dim, fc_dim),
            nn.ReLU(),
            nn.Linear(fc_dim, fc_dim),
            nn.ReLU(),
            nn.Linear(fc_dim, 4)
        )

    def forward(self, x):
        """Output volume density and RGB color (4 dimensions), given a
        positional encoded points sampled from the rays
        """

```

```

x = self.nerf(x)
return x

```

```

def get_minibatches(inputs: torch.Tensor, chunksize: Optional[int] = 1
r"""Takes a huge tensor (ray "bundle") and splits it into a list c
Each element of the list (except possibly the last) has dimension
`chunksize`.
"""
return [inputs[i:i + chunksize] for i in range(0, inputs.shape[0],

```

```

def nerf_step_forward(height, width, focal_length, trans_matrix,
                      near_point, far_point, num_depth_samples_p
                      get_minibatches_function, model):
r"""Perform one iteration of training, which take information of c
training images, and try to predict its rgb values

```

Args:

```

height: height of the image
width: width of the image
focal_length: focal length of the camera
trans_matrix: transformation matrix, which is also the camera pc
near_point: threshold of nearest point
far_point: threshold of farthest point
num_depth_samples_per_ray: number of sampled depth from each ray
get_minibatches_function: function to cut the ray bundles into s
to avoid out-of-memory issue

```

Expected Returns:

```

rgb_predicted: predicted rgb values of the training image
"""

```

YOUR CODE START

```

# Step 1: Generate rays
# Assuming an implementation of `get_rays` function that returns r
ray_origins, ray_directions = get_rays(height, width, focal_length
# print("ray_origins",ray_origins.shape)
# print("ray_directions",ray_directions.shape)
# print(f"ray_origins is stored on: {ray_origins.device}")
# print(f"ray directions is stored on: {ray directions.device}")

```

```

# Step 2: Sample points along each ray
# Assuming an implementation of `sample_points_from_rays` that returns
sampled_points, depth_values = sample_points_from_rays(ray_origins, ray_dirs, num_samples)
# print("sampled_points", sampled_points.shape)
# print("depth_values", depth_values.shape)
# print(f"sampled_points is stored on: {sampled_points.device}")
# print(f"depth_values is stored on: {depth_values.device}")

# Step 3: Apply positional encoding
# Assuming `positional_encoding` expects a flattened list of point coordinates
flattened_sampled_points = sampled_points.reshape(-1, 3) # Flatten
positional_encoded_points = positional_encoding(flattened_sampled_points)
# print("positional_encoded_points", positional_encoded_points.shape)
# print(f"positional_encoded_points is stored on: {positional_encoded_points.device}")

# Step 4: Run the model in batches
# Splitting the points into manageable chunks to avoid OOM
batches = get_minibatches_function(positional_encoded_points, chunk_size)
predictions = []
for batch in batches:
    # print(batch.shape)
    predictions.append(model(batch))

radiance_field_flattened = torch.cat(predictions, dim=0)

##### YOUR CODE START #####

# Step 5: Volume rendering
# Reshape the radiance field to its unflattened shape
radiance_field = radiance_field_flattened.view(height, width, num_channels)
# print(f"radiance_field is stored on: {radiance_field.device}")

# Assuming an implementation of `volume_rendering` that takes the radiance field and ray origins
rgb_predicted = volume_rendering(radiance_field, ray_origins, ray_dirs, num_samples)
# print("rgb_predicted", rgb_predicted.shape)

return rgb_predicted

```



```
return rgb_predicted
```

```
def train(images, poses, hwf, near_point,
          far_point, num_depth_samples_per_ray,
          num_iters, model, History, DEVICE="cuda"):
    r"""Training a tiny nerf model
```

Args:

```
    images: all the images extracted from dataset (including train,
    poses: poses of the camera, which are used as transformation mat
    hwf: [height, width, focal_length]
    near_point: threshold of nearest point
    far_point: threshold of farthest point
    num_depth_samples_per_ray: number of sampled depth from each ray
    num_iters: number of training iterations
    model: predefined tiny NeRF model
    """
```

```
H, W, focal_length = hwf
H = int(H)
W = int(W)
n_train = images.shape[0]
print(H)
```

```
# Optimizer parameters
```

```
lr = 5e-3
```

```
optimizer = torch.optim.Adam(model.parameters(), lr=lr)
```

```
# Seed RNG, for repeatability
```

```
seed = 9458
```

```
torch.manual_seed(seed)
```

```
np.random.seed(seed)
```

```
for _ in tqdm(range(num_iters)):
```

```
    # Randomly pick a training image as the target, get rgb value ar
```

```
    train_idx = np.random.randint(n_train)
```

```
    train_img_rgb = images[train_idx, ..., :3]
```

```
    . . . . .
```

```

train_pose = poses[train_idx]

# Run one iteration of TinyNeRF and get the rendered RGB image.
rgb_predicted = nerf_step_forward(H, W, focal_length,
                                   train_pose, near_point,
                                   far_point, num_depth_samples,
                                   get_minibatches, model)

train_img_rgb_tensor = torch.from_numpy(train_img_rgb)
train_img_rgb = train_img_rgb_tensor.to(DEVICE).to(dtype=torch.float32)

# Compute mean-squared error between the predicted and target image.
loss = torch.nn.functional.mse_loss(rgb_predicted, train_img_rgb)
loss.backward()
optimizer.step()
optimizer.zero_grad()
History.append(loss.cpu().detach().numpy())

if _ % 200 == 0:
    with torch.no_grad():
        plt.figure(figsize=(10, 4))

        plt.subplot(121)
        plt.imshow(train_img_rgb_tensor.cpu().numpy())
        plt.title(f"Target Image at Iteration {_}")

        plt.subplot(122)
        plt.imshow(rgb_predicted.detach().cpu().numpy())
        plt.title(f"Predicted Image at Iteration {_}")

        plt.show()

        print('Loss at ' + str(_) + " iterations: ", loss.cpu().detach().numpy())

print('Finish training')

```

```

if __name__ == "__main__":
    images, poses, hwf=load_colmap_data()
    #print(type(images[0,:,:,:]))

    #images = data["images"]
    #im_shape = images.shape
    #(num_images, H, W, _) = images.shape
    poses= np.array(data["poses"])
    print (H)
    hwf = np.array([H, W, focal])


    # If plotting or processing images, ensure conversion is safe
    plt.imshow(images[np.random.randint(low=0, high=num_images)])
    plt.show()

    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    #device = 'cpu' #Due to OOM
    near_point=2.
    far_point=6.
    num_depth_samples_per_ray = 96
    num_iters = 2000
    model = TinyNeRF(69)
    model.to(device)

    History=[]
    train(images, poses, hwf, near_point,
          far_point, num_depth_samples_per_ray,
          num_iters, model, History,DEVICE=device)

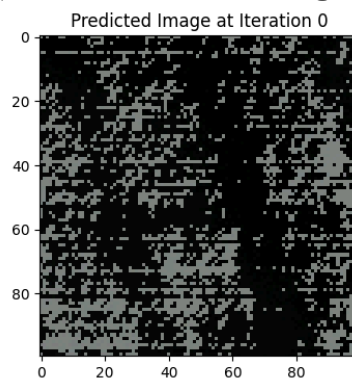

```

100





100

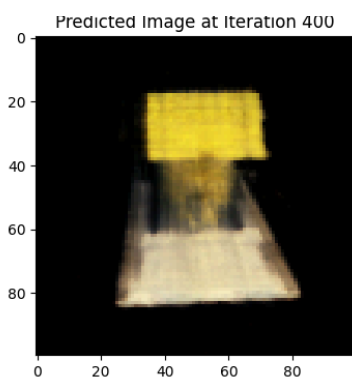
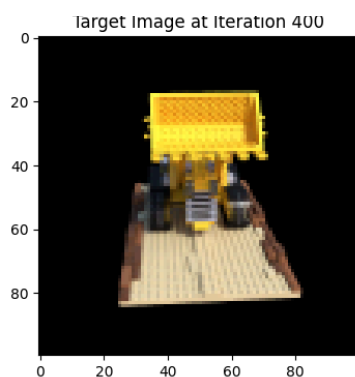
```
0%|          | 0/2000 [00:00<?, ?it/s]  
return _VF.meshgrid(tensors, **kwargs)
```



```
0%|          | 1/2000 [00:08<4:49:22,  
10%|█        | 200/2000 [21:17<3:04:40
```



```
10%|██       | 201/2000 [21:24<3:12:59  
20%|████      | 400/2000 [42:24<3:01:35
```

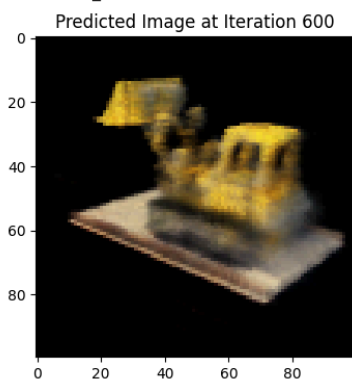
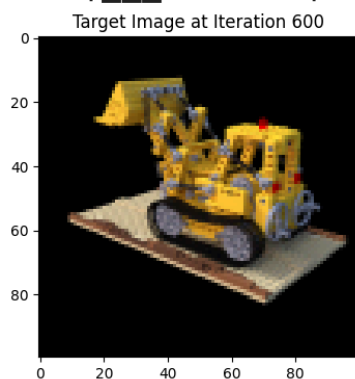


20% |

401/2000 [42:30<2:58:43

30% |

600/2000 [1:04:52<2:41:

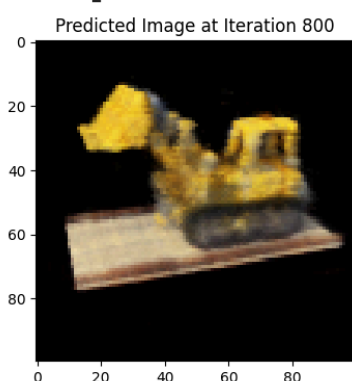



30% |

601/2000 [1:04:59<2:41:

40% |

800/2000 [1:26:38<2:07:

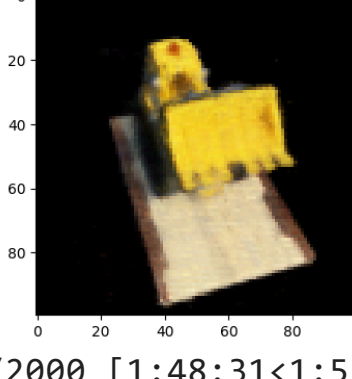
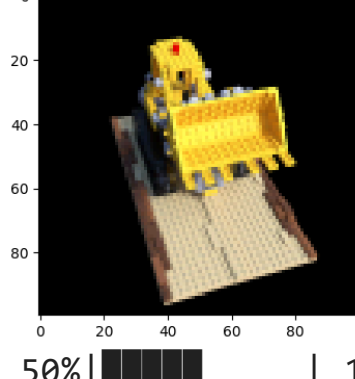


40% |

801/2000 [1:26:44<2:06:

50% |

1000/2000 [1:48:24<1:54





50% |

1001/2000 [1:48:31<1:53

60% |

1200/2000 [2:11:22<1:33





60%



1201/2000 [2:11:30<1:34


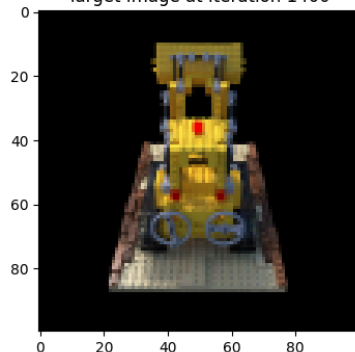
70%



1400/2000 [2:33:17<1:07

Target Image at Iteration 1400

Predicted Image at Iteration 1400



70%



1401/2000 [2:33:24<1:07



80%



1600/2000 [2:55:18<44:0

Target Image at Iteration 1600

Predicted Image at Iteration 1600



80%



1601/2000 [2:55:24<43:0



90%



1800/2000 [3:17:38<22:1

Target Image at Iteration 1800

Predicted Image at Iteration 1800



90%



1801/2000 [3:17:44<21:3

100%



2000/2000 [3:40:06<00:0


```
nlt.plot(History)
```

```
plt.plot(losses, 'b-');  
plt.grid()  
plt.xlabel("Epoch")  
plt.ylabel("Loss")  
plt.title("ESE 650 Zhanqian Nerf Training Process")  
plt.show()
```

