
ESE 650 HW 2

Zhanqian Wu
University of Pennsylvania
Philadelphia, PA 19104
{Zhanqian}@seas.upenn.edu
Coauthor Bowen Jiang bwjiang@seas.upenn.edu

1 Problem 1

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4
5 def generate_transition_matrix(env_shape, obstacles_list, goal_idx):
6     # The environment's shape is a square grid of size 'n'
7     n = env_shape[0]
8
9     # Initialize the transition matrix with zeros
10    # The matrix is 5-dimensional: state (i, j), action, resulting state (new_i, new_j)
11    # Actions are encoded as 0: left, 1: right, 2: up, 3: down
12    T = np.zeros((n, n, 4, n, n))
13
14    # Iterate over all cells in the grid to set transition probabilities
15    for i in range(n):
16        for j in range(n):
17            # If the cell is the goal state, any action results in staying in the goal with
18            # probability 1
19            if (i, j) == goal_idx:
20                T[i, j, :, i, j] = 1
21                continue # Skip further processing for the goal state
22
23            # If the cell is an obstacle, skip it as no action is applicable
24            if (i, j) in obstacles_list:
25                continue
26
27            # Define transitions based on the control action taken by the robot
28            # For each direction, there's a primary movement direction with p=0.7
29            # and secondary movements (including staying in place) with p=0.1
30
31            # For action 0 (left)
32            T[i, j, 0, i, np.clip(j-1, 0, n-1)] += 0.7
33            T[i, j, 0, np.clip(i-1, 0, n-1), j] += 0.1
34            T[i, j, 0, np.clip(i+1, 0, n-1), j] += 0.1
35            T[i, j, 0, i, j] += 0.1
36
37            # For action 1 (right)
38            T[i, j, 1, i, np.clip(j+1, 0, n-1)] += 0.7
39            T[i, j, 1, np.clip(i-1, 0, n-1), j] += 0.1
40            T[i, j, 1, np.clip(i+1, 0, n-1), j] += 0.1
41            T[i, j, 1, i, j] += 0.1
```

```

42         # For action 2 (up)
43         T[i, j, 2, np.clip(i-1, 0, n-1), j] += 0.7
44         T[i, j, 2, i, np.clip(j+1, 0, n-1)] += 0.1
45         T[i, j, 2, i, np.clip(j-1, 0, n-1)] += 0.1
46         T[i, j, 2, i, j] += 0.1
47
48         # For action 3 (down)
49         T[i, j, 3, np.clip(i+1, 0, n-1), j] += 0.7
50         T[i, j, 3, i, np.clip(j+1, 0, n-1)] += 0.1
51         T[i, j, 3, i, np.clip(j-1, 0, n-1)] += 0.1
52         T[i, j, 3, i, j] += 0.1
53
54     return T
55
56
57 def generate_state_map(env_shape):
58     """
59     Generates a map of the environment with obstacles and points of interest.
60
61     Args:
62         env_shape (tuple): The dimensions of the environment (height, width).
63
64     Returns:
65         np.ndarray: A 2D array representing the environment where
66                     0 = free space,
67                     1 = obstacle,
68                     2 = goal,
69                     3 = initial position,
70                     4 = point of interest.
71     """
72     # Initialize the environment with free spaces
73     state_map = np.zeros(env_shape)
74
75     # Set the boundaries of the environment as obstacles
76     state_map[:, 0] = state_map[:, -1] = state_map[0, :] = state_map[-1, :] = 1
77
78     # Add specific obstacles within the environment
79     state_map[2, 3:7] = 1
80     state_map[4:8, 4] = 1
81     state_map[7, 5] = 1
82     state_map[4:6, 7] = 1
83
84     # Flip the map if needed to match the desired orientation
85     # state_map = np.flip(state_map, axis=1)
86     # state_map = np.fliplr(state_map)
87
88     # Mark specific points of interest
89     state_map[8, 8] = 2 # Goal position
90     state_map[8, 1] = 3 # Initial position
91     state_map[3, 3] = 4 # Another point of interest
92
93     return state_map
94
95
96 def generate_Q_map(states, env_shape, reward):
97     """
98     Generate the Q-value map for an environment.
99
100    Args:
101        states (np.ndarray): A 2D array where cells are marked with 0 for free space,
102                             1 for obstacles, and 2 for the goal.
103        env_shape (tuple): Shape of the environment as (height, width).
104        reward (float): The reward for reaching the goal or hitting an obstacle.
105
106    Returns:

```

```

107     np.ndarray: A 3D array representing the Q-values for each action at each cell.
108         The shape of the array is (height, width, 4), corresponding to
109         the environment dimensions and four possible actions.
110     """
111
112     # Identify the coordinates of obstacles and the goal in the grid
113     obstacles = np.where(states == 1)
114     goal_idx = np.where(states == 2)
115
116     # Initialize the Q-value map with -1 for all states and actions
117     # This encourages exploration by giving a slightly negative value to unvisited states
118     Q = -1 * np.ones((env_shape[1], env_shape[0], 4)) # Notice the corrected order of dimensions
119
120     # Assign a negative reward to all actions leading to obstacle states
121     # This penalizes hitting obstacles
122     Q[obstacles[0], obstacles[1], :] = -reward
123
124     # Assign a positive reward to all actions leading to the goal state
125     # This incentivizes reaching the goal
126     Q[goal_idx[0], goal_idx[1], :] = reward
127
128     return Q
129
130
131
132 def policy_evaluation(T, Q, J_init, u):
133     """
134     Evaluates a policy to estimate the state-value function for each state.
135
136     J[k, i, j]: This represents the estimated value (utility) of being in state (i, j) at iteration
137     k under a specific policy u. The value function J estimates the total expected rewards from
138     being in a particular state and following a certain policy thereafter.
139
140     Q[i, j, u[i, j]]: This is the immediate reward (or the action-value) of taking action u[i, j] (
141     the action recommended by the policy at state (i, j)) plus the expected future rewards. The Q
142     matrix stores the action-value function, which gives the quality of each action at each state.
143
144     gamma: This is the discount factor (denoted as gamma). It represents the difference in
145     importance
146
147     J[k-1].flatten().T: This term represents the value function from the previous iteration (k-1)
148
149     Args:
150         T (np.ndarray): The transition probabilities matrix of shape (n, n, 4, n, n).
151         Q (np.ndarray): The action-value function matrix of shape (env_height, env_width, actions).
152         J_init (np.ndarray): Initial state-value function of shape (env_height, env_width).
153         u (np.ndarray): Policy matrix indicating the action for each state.
154
155     Returns:
156         np.ndarray: Estimated state-value function after policy evaluation.
157     """
158     iter = 300
159     J = np.zeros((iter, env_shape[0], env_shape[1]))
160     J[0] = J_init
161     for k in range(1, iter):
162         for i in range(10):
163             for j in range(10):
164                 J[k,i,j] = Q[i,j,u[i,j]] + gamma*(T[i,j,u[i,j]].flatten() @ J[k-1].flatten().T)
165
166     return J[-1]
167
168 def policy_improvement(T, Q, J_new, gamma, env_shape):

```

```

167 """
168 Performs policy improvement by finding an improved policy based on the updated state-value
    function.
169
170 Args:
171     T (np.ndarray): Transition probabilities matrix of shape (env_height, env_width, actions,
    env_height, env_width).
172     Q (np.ndarray): Action-value function matrix of shape (env_height, env_width, actions).
173     J_new (np.ndarray): Updated state-value function of shape (env_height, env_width) from
    policy evaluation.
174     gamma (float): Discount factor for future rewards.
175     env_shape (tuple): Shape of the environment (height, width).
176
177 Returns:
178     np.ndarray: Improved policy matrix indicating the best action for each state.
179 """
180 # Define a lambda function to find the action that maximizes the expected utility for each
    state
181 get_best_action = lambda action_values: np.argmax(action_values)
182
183 # Initialize the improved policy matrix with zeros
184 u_k = np.zeros(env_shape)
185
186 # Iterate over all states in the environment
187 for i in range(env_shape[0]):
188     for j in range(env_shape[1]):
189         # Extract Q-values and transition probabilities for the current state
190         Q_element = Q[i, j]
191         T_element = T[i, j].reshape(4, -1)
192
193         # Compute the expected utility for each action and select the best action
194         u_k[i, j] = get_best_action(Q_element + (gamma * T_element @ J_new.reshape((-1, 1))).
    reshape(4))
195
196 # Return the improved policy
197 return u_k
198
199 def visualize(J_new, u, obstacles_list, goal_idx_list, iteration_number):
200     """
201     Visualizes the policy and value function of a grid.
202
203     Args:
204         J_new (np.ndarray): The value function to be visualized.
205         u (np.ndarray): The policy matrix, with actions for each cell.
206         obstacles_list (list of tuples): Coordinates of obstacles in the grid.
207         goal_idx_list (list of tuples): Coordinates of goal(s) in the grid.
208     """
209
210     # Set up the figure and axes for the visualization
211     fig, ax = plt.subplots(figsize=(5, 5))
212     ax.set_xticks(np.arange(0.5, 10.5, 1))
213     ax.set_yticks(np.arange(0.5, 10.5, 1))
214
215     # Use a heatmap to visualize the value function
216     cmap = plt.get_cmap('bwr')
217     im = plt.imshow(J_new, cmap=cmap)
218     plt.imshow(J_new, cmap=cmap)
219     plt.grid(color='gray', linestyle='--', linewidth=0.5)
220
221     # Draw arrows to represent the policy at each cell
222     for i in range(10):
223         for j in range(10):
224             if (i, j) not in obstacles_list and (i, j) not in goal_idx_list:
225                 dx, dy = 0, 0
226                 if u[i, j] == 0: # Left

```

```

227         dx = -0.25
228         elif u[i, j] == 1: # Right
229             dx = 0.25
230         elif u[i, j] == 2: # Up
231             dy = -0.25
232         elif u[i, j] == 3: # Down
233             dy = 0.25
234         ax.arrow(j, i, dx, dy, head_width=0.1, head_length=0.1, fc='red', ec='red') #
Change arrow color here
235
236     fig.colorbar(im, ax=ax)
237     plt.title(f"Actions taken at the {iteration_number} iteration(s)")
238     plt.show()
239
240
241 # Initial setup
242 env_shape = (10, 10)
243 reward = 10
244 gamma = 0.9
245
246 # Generate environment states, transition matrix, and initial Q-values
247 states = generate_state_map(env_shape)
248 obstacles_list = list(zip(*np.where(states == 1)))
249 goal_idx_list = list(zip(*np.where(states == 2)))
250 T = generate_transition_matrix(env_shape, obstacles_list, goal_idx_list[0])
251 Q = generate_Q_map(states, env_shape, reward)
252
253 # Initialize policy evaluation and improvement
254 num_policy_iter = 4
255 J = np.zeros(env_shape)
256 u = np.ones((num_policy_iter+1, env_shape[0], env_shape[1]), dtype=int)
257
258 # Iterate over policy evaluation and improvement
259 for k in range(num_policy_iter):
260     J_new = policy_evaluation(T, Q, J, u[k])
261     J = J_new
262     u[k+1] = policy_improvement(T, Q, J_new, gamma, env_shape)
263     visualize(J_new, u[k+1], obstacles_list, goal_idx_list, k)

```

Code Listing 1: Policy Iteration Code

1.1 (b) Initialize policy iteration with a feedback control




Figure 1: value function $J_{\pi}(0)(x)$ as a heatmap in the above picture

1.2 (c) Execute the policy iteration algorithm

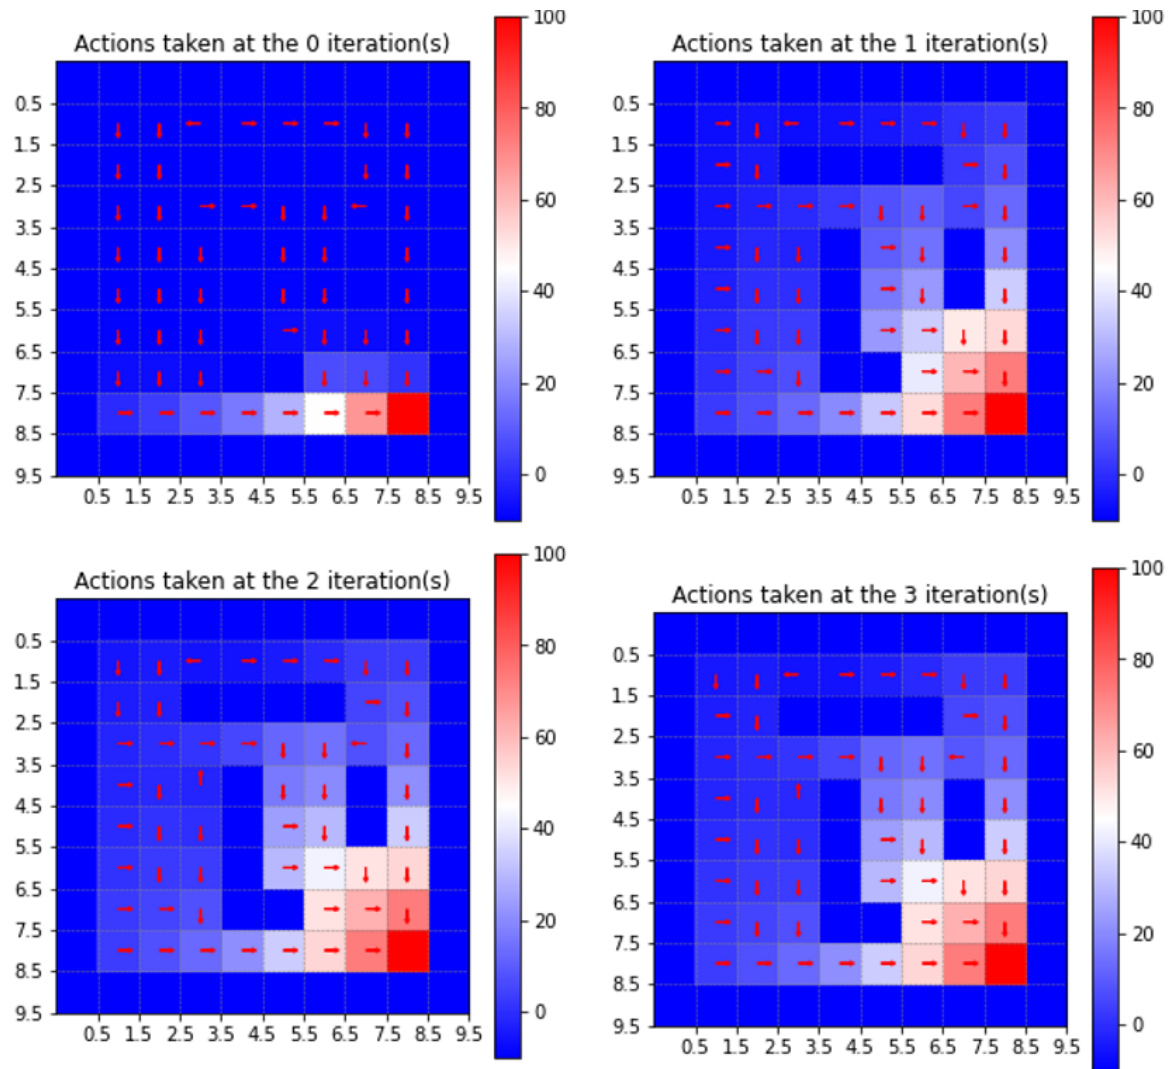


Figure 2: Plot the feedback control for the first four iterations

2 Problem 2

```
1 import os, sys, pickle, math
2 from copy import deepcopy
3
4 from scipy import io
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 from load_data import load_lidar_data, load_joint_data, joint_name_to_index
9 from utils import *
10
11 import logging
12
13 logger = logging.getLogger()
14 logger.setLevel(os.environ.get("LOGLEVEL", "INFO"))
15
16
17 class map_t:
18     """
19     This will maintain the occupancy grid and log_odds.
20     You do not need to change anything in the initialization
21     """
22
23     def __init__(s, resolution=0.05):
24         s.resolution = resolution
25         s.xmin, s.xmax = -20, 20
26         s.ymin, s.ymax = -20, 20
27         s.szx = int(np.ceil((s.xmax - s.xmin) / s.resolution + 1))
28         s.szy = int(np.ceil((s.ymax - s.ymin) / s.resolution + 1))
29
30         # binarized map and log-odds
31         s.cells = np.zeros((s.szx, s.szy), dtype=np.int8) # initialize the map as empty
32         s.log_odds = np.zeros(s.cells.shape, dtype=np.float64)
33
34         s.log_odds_max = 5e6
35         # number of observations received yet for each cell
36         s.num_obs_per_cell = np.zeros(s.cells.shape, dtype=np.uint64)
37
38         # we call a cell occupied if the probability of
39         # occupancy  $P(m_i | \dots)$  is  $\geq$  occupied_prob_thresh
40         s.occupied_prob_thresh = 0.6
41         s.log_odds_thresh = np.log(s.occupied_prob_thresh / (1 - s.occupied_prob_thresh))
42
43     def grid_cell_from_xy(s, x, y):
44         """
45         x and y are 1-dimensional arrays, compute the cell indices in the map corresponding to
46         these (x,y) locations.
47         You should return an array of shape (2 x len(x)).
48         Be careful to handle instances when x/y go outside the map bounds,
49         you can use np.clip to handle these situations.
50         """
51         ##### TODO: Checked
52         x_indices = np.clip((x - s.xmin) / s.resolution, 0, s.szx - 1).astype(int)
53         y_indices = np.clip((y - s.ymin) / s.resolution, 0, s.szy - 1).astype(int)
54
55         # Stack the x_indices and y_indices vertically to create a 2D array where each column
56         # represents
57         return np.vstack((x_indices, y_indices))
58
59 class slam_t:
60     """
61     s is the same as s. In Python it does not really matter
62     """
```



```

61 what we call s, s is shorter. As a general comment, (I believe)
62 you will have fewer bugs while writing scientific code if you
63 use the same/similar variable names as those in the mathematical equations.
64 """
65
66 def __init__(s, resolution=0.05, Q=1e-3 * np.eye(3), resampling_threshold=0.3):
67
68     s.init_sensor_model()
69
70     # dynamics noise for the state (x,y,yaw)
71     s.Q = Q
72     # s.Q = 1e-8 * np.eye(3)
73
74     # we resample particles if the effective number of particles
75     # falls below s.resampling_threshold*num_particles
76     s.resampling_threshold = resampling_threshold
77
78     # initialize the map
79     s.map = map_t(resolution)
80
81 def read_data(s, src_dir, idx=0, split='train'):
82     """
83     src_dir: location of the "data" directory
84     """
85     logging.info('> Reading data')
86     s.idx = idx
87     s.lidar = load_lidar_data(os.path.join(src_dir, 'data/%s/%s_lidar%d' % (split, split, idx))
88 )
89     s.joint = load_joint_data(os.path.join(src_dir, 'data/%s/%s_joint%d' % (split, split, idx))
90 )
91
92     # finds the closets idx in the joint timestamp array such that the timestamp
93     # at that idx is t
94     s.find_joint_t_idx_from_lidar = lambda t: np.argmin(np.abs(s.joint['t'] - t))
95
96 def init_sensor_model(s):
97     # lidar height from the ground in meters
98     s.head_height = 0.93 + 0.33
99     s.lidar_height = 0.15
100
101     s.lidar_dmin = 1e-3
102     s.lidar_dmax = 30
103     s.lidar_angular_resolution = 0.25 # degrees
104     s.lidar_angles = np.arange(-135, 135 + s.lidar_angular_resolution,
105                               s.lidar_angular_resolution) * np.pi / 180.0
106
107     # Sensor model
108     s.lidar_log_odds_occ = np.log(9)
109     s.lidar_log_odds_free = np.log(1 / 9.)
110
111 def init_particles(s, n=100, p=None, w=None, t0=0):
112     """
113     n: number of particles
114     p: xy yaw locations of particles (3xn array)
115     w: weights (array of length n)
116     """
117     s.n = n
118     s.p = deepcopy(p) if p is not None else np.zeros((3, s.n), dtype=np.float64)
119     s.w = deepcopy(w) if w is not None else np.ones(n) / float(s.n) # 1/n
120
121 @staticmethod
122 def stratified_resampling(p, w):
123     """
124     resampling step of the particle filter, takes p = 3 x n array of
125     particles with w = 1 x n array of weights and returns new particle
126     locations (number of particles n remains the same) and their weights

```

```

124 Parameters:
125 - p: (3 x n) numpy array of particle states.
126 - w: (1 x n) numpy array of particle weights.
127
128 Returns:
129 - Tuple of (resampled_particles, uniform_weights):
130   - resampled_particles is a (3 x n) numpy array after resampling.
131   - uniform_weights is a (1 x n) numpy array of equal weights for all particles.
132 """
133 ##### TODO: Checked
134 n = len(w) # Total number of particles
135 # Adjust weights and repeat particles based on adjusted weights
136 adjusted_weights = (w * n * 10).astype(int)
137 repeated_particles = np.repeat(p, adjusted_weights, axis=1)
138
139 # Select a subset of particles randomly
140 indexes = np.random.choice(repeated_particles.shape[1], n, replace=False)
141 resampled_particles = repeated_particles[:, indexes]
142
143 # Assign equal weight to all resampled particles
144 uniform_weights = np.ones(n) / n
145
146 return resampled_particles, uniform_weights
147
148
149 @staticmethod
150 def log_sum_exp(w):
151     return w.max() + np.log(np.exp(w - w.max()).sum())
152
153
154 def rays2world(s, p, d, head_angle=0, neck_angle=0, angles=None):
155     """
156     p: p is the pose of the particle (x,y,yaw), a 3x1 array describing the robot position and
157     orientation
158     d: an array that stores the distance along the ray of the lidar for each ray
159     the length of d has to be equal to that of angles, this is s.lidar[t]['scan']
160     head_angle: the angle of the head in the body frame, usually 0, need to be in radians
161     neck_angle: the angle of the neck in the body frame, usually 0, need to be in radians
162     angles: angle of each ray in the body frame in radians
163             (usually be simply s.lidar_angles for the different lidar rays)
164
165     Return an array (2 x num_rays) which are the (x,y) locations of the end point of each ray
166     in world coordinates
167     """
168     # Filter valid LiDAR points based on distance constraints
169     in_range = np.logical_and(d >= s.lidar_dmin, d <= s.lidar_dmax)
170     d_filtered = d[in_range]
171     angle_filtered = angles[in_range]
172
173     # Transform distances to LiDAR frame points (2D to 3D)
174     lidar_pts = np.vstack((d_filtered * np.cos(angle_filtered), d_filtered * np.sin(
175     angle_filtered), np.zeros(d_filtered.size)))
176
177     # Transformation from LiDAR to body frame
178     lidar_to_body_tf = euler_to_se3(0, head_angle, neck_angle, np.array([0, 0, s.lidar_height])
179 )
180     body_pts_4d = lidar_to_body_tf @ make_homogeneous_coords_3d(lidar_pts) # Transform to 4D
181     for matrix multiplication
182
183     # Transform from body frame to world frame
184     body_to_world_tf = euler_to_se3(0, 0, p[2, 0], np.array([p[0, 0], p[1, 0], s.head_height]))
185     world_pts_4d = body_to_world_tf @ body_pts_4d # Apply transformation
186
187     # Normalize and return 2D world frame points
188     world_pts_2d = world_pts_4d[:3] / world_pts_4d[3]

```

```

184     return world_pts_2d[:2]
185
186 def get_control(s, t):
187     """
188     Use the pose at time t and t-1 to calculate what control the robot could have taken
189     at time t-1 at state (x,y,th)_{t-1} to come to the current state (x,y,th)_t. We will
190     assume that this is the same control that the robot will take in the function dynamics_step
191     below at time t, to go to time t-1. need to use the smart_minus_2d function to get the
    difference of the two poses and we will simply set this to be the control (delta x, delta y,
    delta theta)
    """
192
193     """
194
195     Parameters:
196     - t: The current time step (index) in the LIDAR data sequence.
197
198     Returns:
199     - control: The computed control signal as a difference in x, y coordinates,
200               and heading angle (theta),
201               indicating how to adjust the pose from time step t-1 to t.
202     """
203
204     if t == 0:
205         return np.zeros(3)
206
207     ##### TODO: Checked
208
209     # Compute the difference in pose between time t and t-1
210     # Extract the previous and current poses from LIDAR data using the given time step t.
211     previous_pose = s.lidar[t - 1]['xyth']
212     current_pose = s.lidar[t]['xyth']
213
214     # Compute the control signal as the difference between current and previous poses.
215     return smart_minus_2d(current_pose, previous_pose)
216
217 def dynamics_step(s, t):
218     """
219     Compute the control using get_control and perform that control on each particle to get the
220     updated locations of the particles in the particle filter, remember to add noise using the
    smart_plus_2d function to each particle
    """
221
222     """
223
224     Parameters:
225     - control: The control signal to be applied, typically the difference in pose.
226     """
227     ##### TODO: Checked
228     control = s.get_control(t)
229
230     # Generate noise for all particles at once
231     noise = np.random.multivariate_normal(np.zeros(s.Q.shape[0]), s.Q, s.n)
232
233     # Apply the noisy control to each particle
234     for i in range(s.n):
235         noisy_control = control + noise[i]
236         s.p[:, i] = smart_plus_2d(s.p[:, i].copy(), noisy_control)
237
238 @staticmethod
239 def update_weights(w, obs_logp):
240     """
241     Given the observation log-probability and the weights of particles w, calculate the
242     new weights as discussed in the writeup. Make sure that the new weights are normalized
243     """
244     # Parse the observation log-probability

```

```

245     w = obs_logp + np.log(w)
246     w -= slam_t.log_sum_exp(w)
247     w = np.exp(w)
248     return w
249
250 def observation_step(s, t):
251     """
252     This function does the following things
253     1. updates the particles using the LiDAR observations
254     2. updates map.log_odds and map.cells using occupied cells as shown by the LiDAR data
255
256     Some notes about how to implement this.
257     1. As mentioned in the writeup, for each particle
258         (a) First find the head, neck angle at t (this is the same for every particle)
259         (b) Project lidar scan into the world frame (different for different particles)
260         (c) Calculate which cells are obstacles according to this particle for this scan,
261             calculate the observation log-probability
262     2. Update the particle weights using observation log-probability
263     3. Find the particle with the largest weight, and use its occupied cells to update the
264     map.log_odds and map.cells.
265     You should ensure that map.cells is recalculated at each iteration (it is simply the
266     binarized version of log_odds). map.log_odds is of course maintained across iterations.
267     """
268     ##### TODO: checked
269     # Extract head and neck angles
270     idx = s.find_joint_t_idx_from_lidar(s.lidar[t]['t'])
271     angle_neck = s.joint['head_angles'][0, idx]
272     angle_head = s.joint['head_angles'][1, idx]
273
274     # Initialize observation probabilities
275     log_prob_obs = np.zeros(s.n)
276
277     for i in range(s.n):
278         # Project lidar scan ---> world frame
279         p = s.p[:, i].reshape((3, 1))
280         world_frame_points = s.rays2world(p, s.lidar[t]['scan'], angle_head, angle_neck, s.
281         lidar_angles)
282
283         # grid cell indices of the occupied cells && observation log-probability
284         occupied_cells = s.map.grid_cell_from_xy(world_frame_points[0], world_frame_points[1])
285         log_prob_obs[i] = np.sum(s.map.log_odds[occupied_cells[0], occupied_cells[1]])
286
287     # Update particle weights and estimate the pose from the best particle
288     s.w = s.update_weights(s.w, log_prob_obs)
289     best_idx = np.argmax(s.w)
290     s.estimated_pose = s.p[:, best_idx]
291
292     # Update the map based on the best particle's observation
293     best_particle_world = s.rays2world(s.estimated_pose.reshape((3, 1)), s.lidar[t]['scan'],
294     angle_head, angle_neck, s.lidar_angles)
295     occupied_x, occupied_y = s.map.grid_cell_from_xy(best_particle_world[0],
296     best_particle_world[1])
297
298     # Compute and update free cells from best particle to observed obstacles
299     limits_x, limits_y = s.calculate_free_space(s.estimated_pose, occupied_x, occupied_y)
300     s.update_map(occupied_x, occupied_y, limits_x, limits_y)
301
302     # Resample particles
303     s.resample_particles()
304
305 def calculate_free_space(s, pose, occupied_x, occupied_y):
306     """
307     Calculate free space coordinates based on the current pose and observed occupied cells.
308     """
309     # Compute limits based on lidar maximum distance and current pose

```

```

305     limit_x = np.array([pose[0] - s.lidar_dmax / 2, pose[0] + s.lidar_dmax / 2, pose[0]])
306     limit_y = np.array([pose[1] - s.lidar_dmax / 2, pose[1] + s.lidar_dmax / 2, pose[1]])
307     limit_grid_x, limit_grid_y = s.map.grid_cell_from_xy(limit_x, limit_y)
308
309     # Determine free cells
310     free_x = np.linspace(limit_grid_x[2], occupied_x, endpoint=False).astype(int).flatten()
311     free_y = np.linspace(limit_grid_y[2], occupied_y, endpoint=False).astype(int).flatten()
312
313     return free_x, free_y
314
315
316 def update_map(s, occupied_x, occupied_y, free_x, free_y):
317     """
318     Update SLAM map log-odds values based on observed and free cells, then binarize the map.
319     """
320     # Update log-odds for occupied and free cells
321     s.map.log_odds[occupied_x, occupied_y] += s.lidar_log_odds_occ
322     s.map.log_odds[free_x, free_y] += s.lidar_log_odds_free
323     s.map.log_odds = np.clip(s.map.log_odds, -s.map.log_odds_max, s.map.log_odds_max)
324
325     # Binarize the map based on log-odds threshold
326     s.map.cells = (s.map.log_odds > s.map.log_odds_thresh).astype(int)
327
328 def resample_particles(s):
329     """
330     Resampling is a (necessary) but problematic step which introduces a lot of variance in the
331     particles.
332     We should resample only if the effective number of particles falls below
333     a certain threshold (resampling_threshold).
334     A good heuristic to calculate the effective particles is  $1/(\sum_i w_i^2)$  where  $w_i$  are the
335     weights
336     of the particles, if this number of close to n, then all particles have about equal
337     weights,
338     and we do not need to resample
339     """
340     e = 1 / np.sum(s.w ** 2)
341     logging.debug('> Effective number of particles: {}'.format(e))
342     if e / s.n < s.resampling_threshold:
343         s.p, s.w = s.stratified_resampling(s.p, s.w)
344         logging.debug('> Resampling')

```

Code Listing 2: SLAM Code

```

1
2 # Pratik Chaudhari (pratikac@seas.upenn.edu)
3 import os
4 os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"
5
6 import click, tqdm, random
7 import numpy as np
8 from slam import *
9
10 def run_dynamics_step(src_dir, log_dir, idx, split, t0=0, draw_fig=False):
11     """
12     This function is for you to test your dynamics update step. It will create
13     two figures after you run it. The first one is the robot location trajectory
14     using odometry information obtained from the lidar. The second is the trajectory
15     using the PF with a very small dynamics noise. The two figures should look similar.
16     """
17     slam = slam_t(Q=1e-8*np.eye(3))
18     slam.read_data(src_dir, idx, split)
19
20     # trajectory using odometry (xy and yaw) in the lidar data
21     d = slam.lidar
22     xyth = []
23     for p in d:
24         xyth.append([p['xyth'][0], p['xyth'][1], p['xyth'][2]])
25     xyth = np.array(xyth)
26
27     plt.figure(1); plt.clf();
28     plt.title('Trajectory using onboard odometry')
29     plt.plot(xyth[:,0], xyth[:,1])
30     logging.info('> Saving odometry plot in '+os.path.join(log_dir, 'odometry_%s_%02d.jpg'%(split,
31     idx)))
32     plt.savefig(os.path.join(log_dir, 'odometry_%s_%02d.jpg'%(split, idx)))
33
34     # dynamics propagation using particle filter
35     # n: number of particles, w: weights, p: particles (3 dimensions, n particles)
36     # S covariance of the xyth location
37     # particles are initialized at the first xyth given by the lidar
38     # for checking in this function
39     n = 3
40     w = np.ones(n)/float(n)
41     p = np.zeros((3,n), dtype=np.float64)
42     slam.init_particles(n,p,w)
43     slam.p[:,0] = deepcopy(slam.lidar[0]['xyth'])
44
45     print('> Running prediction')
46     t0 = 0
47     T = len(d)
48     ps = deepcopy(slam.p) # maintains all particles across all time steps
49     plt.figure(2); plt.clf();
50     ax = plt.subplot(111)
51
52     for t in tqdm.tqdm(range(t0+1,T)):
53         slam.dynamics_step(t)
54         ps = np.hstack((ps, slam.p))
55
56         if draw_fig:
57             ax.clear()
58             ax.plot(slam.p[0], slam.p[0], '*r')
59             plt.title('Particles %03d'%t)
60             plt.draw()
61             plt.pause(0.01)
62
63     plt.plot(ps[0], ps[1], '*c')
64     plt.title('Trajectory using PF')

```

```

64 logging.info('> Saving plot in '+os.path.join(log_dir, 'dynamics_only_%s_%02d.jpg'%(split, idx)
65 ))
66 plt.savefig(os.path.join(log_dir, 'dynamics_only_%s_%02d.jpg'%(split, idx)))
67 def run_observation_step(src_dir, log_dir, idx, split, is_online=False):
68     """
69     This function is for you to debug your observation update step
70     It will create three particles np.array([[0.2, 2, 3],[0.4, 2, 5],[0.1, 2.7, 4]])
71     * Note that the particle array has the shape 3 x num_particles so
72     the first particle is at [x=0.2, y=0.4, z=0.1]
73     This function will build the first map and update the 3 particles for one time step.
74     After running this function, you should get that the weight of the second particle is the
75     largest since it is the closest to the origin [0, 0, 0]
76     """
77     slam = slam_t(resolution=0.05)
78     slam.read_data(src_dir, idx, split)
79
80     # t=0 sets up the map using the yaw of the lidar, do not use yaw for
81     # other timestep
82     # initialize the particles at the location of the lidar so that we have some
83     # occupied cells in the map to calculate the observation update in the next step
84     t0 = 0
85     xyth = slam.lidar[t0]['xyth']
86     xyth[2] = slam.lidar[t0]['rpy'][2]
87     logging.debug('> Initializing 1 particle at: {}'.format(xyth))
88     slam.init_particles(n=1,p=xyth.reshape((3,1)),w=np.array([1]))
89
90     slam.observation_step(t=0)
91     logging.info('> Particles\n: {}'.format(slam.p))
92     logging.info('> Weights: {}'.format(slam.w))
93
94     # reinitialize particles, this is the real test
95     logging.info('\n')
96     n = 3
97     w = np.ones(n)/float(n)
98     p = np.array([[2, 0.2, 3],[2, 0.4, 5],[2.7, 0.1, 4]])
99     slam.init_particles(n, p, w)
100
101     slam.observation_step(t=1)
102     logging.info('> Particles\n: {}'.format(slam.p))
103     logging.info('> Weights: {}'.format(slam.w))
104
105 def run_slam(src_dir, log_dir, idx, split):
106     """
107     This function runs slam. We will initialize the slam just like the observation_step
108     before taking dynamics and observation updates one by one. You should initialize
109     the slam with n=100 particles, you will also have to change the dynamics noise to
110     be something larger than the very small value we picked in run_dynamics_step function
111     above.
112     """
113     slam = slam_t(resolution=0.05, Q=np.diag([2e-4,2e-4,1e-4]))
114     slam.read_data(src_dir, idx, split)
115     T = len(slam.lidar)
116
117     # again initialize the map to enable calculation of the observation logp in
118     # future steps, this time we want to be more careful and initialize with the
119     # correct lidar scan. First find the time t0 around which we have both LiDAR
120     # data and joint data
121
122     #### TODO: Checked
123
124     # initialize the occupancy grid using one particle and calling the observation_step
125     # function
126
127     # Args:

```

```

127 # - src_dir: Source directory containing the dataset.
128 # - log_dir: Directory to save the output map visualization.
129 # - idx: Index of the dataset to process.
130 # - split: Dataset split (e.g., 'train', 'test') to use.
131
132 ##### TODO: Checked
133
134 # Initialize SLAM with specific noise parameters and resolution
135 slam = slam_t(resolution=0.05, Q=np.diag([2e-4, 2e-4, 1e-4]))
136 slam.read_data(src_dir, idx, split)
137 total_steps = len(slam.lidar)
138
139 # Initialize occupancy grid with one particle at the first lidar position
140 init_pose = np.array(slam.lidar[0]['xyth'])
141 init_pose[2] = slam.lidar[0]['rpy'][2] # Use the yaw from 'rpy'
142 slam.init_particles(n=1, p=init_pose.reshape((3, 1)), w=np.array([1.0]))
143 slam.observation_step(t=0)
144
145 # Log initialization info
146 logging.info(f'> Initializing particles at first timestamp: {slam.lidar[0]["t"]}')
147
148 # Re-initialize SLAM with 100 particles
149 slam.init_particles(n=100)
150 slam.dynamics_step(0)
151
152 # Process dynamics and observations for each time step
153 particle_positions = []
154 for t in tqdm.tqdm(range(1, total_steps)):
155     slam.dynamics_step(t)
156     slam.observation_step(t)
157     particle_positions.append(slam.estimated_pose)
158
159 # Convert list of particle positions to a numpy array
160 particle_positions = np.array(particle_positions)
161
162 # Plot and save the map and particle trajectories
163 fig, ax = plt.subplots(figsize=(10, 10))
164 occupied_x, occupied_y = np.where(slam.map.cells == 1) # Plot occupied cells
165 ax.plot(occupied_x, occupied_y, 'sk', markersize=1, label='Occupied') # Plot particles
166 particle_x, particle_y = slam.map.grid_cell_from_xy(particle_positions[:, 0],
167     particle_positions[:, 1])
168 ax.plot(particle_x, particle_y, '.r', markersize=5, label='Particles')
169 ax.grid(True)
170 ax.legend()
171 ax.set(xlim=(0, slam.map.szx), ylim=(0, slam.map.szy), title=f'Map {idx}')
172 ax.set_xlabel('X Coordinate')
173 ax.set_ylabel('Y Coordinate')
174 plt.show()
175 fig.savefig(os.path.join(log_dir, f'Map{idx}_{split}.png'))
176 plt.close(fig)
177
178 @click.command()
179 @click.option('--src_dir', default='./', help='data directory', type=str)
180 @click.option('--log_dir', default='logs', help='directory to save logs', type=str)
181 @click.option('--idx', default='0', help='dataset number', type=int)
182 @click.option('--split', default='train', help='train/test split', type=str)
183 @click.option('--mode', default='slam',
184     help='choices: dynamics OR observation OR slam', type=str)
185 def main(src_dir, log_dir, idx, split, mode):
186     # Run python main.py --help to see how to provide command line arguments
187
188     if not mode in ['slam', 'dynamics', 'observation']:
189         raise ValueError('Unknown argument --mode %s'%mode)
190     sys.exit(1)

```



```

191 np.random.seed(42)
192 random.seed(42)
193
194
195 if mode == 'dynamics':
196     run_dynamics_step(src_dir, log_dir, idx, split)
197     sys.exit(0)
198 elif mode == 'observation':
199     run_observation_step(src_dir, log_dir, idx, split)
200     sys.exit(0)
201 else:
202     p = run_slam(src_dir, log_dir, idx, split)
203     return p
204
205 if __name__ == '__main__':
206     src_dir = './'
207     log_dir = 'logs'
208     idx = 3
209     split = 'train'
210
211     run_dynamics_step(src_dir, log_dir, idx, split)
212     run_observation_step(src_dir, log_dir, idx, split)
213     run_dynamics_step(src_dir, log_dir, idx, split)
214     run_slam(src_dir, log_dir, idx, split)
215     main()

```

Code Listing 3: Simultaneous Localization and Mapping (SLAM) with a particle filter main code

2.1 (c) dynamics step

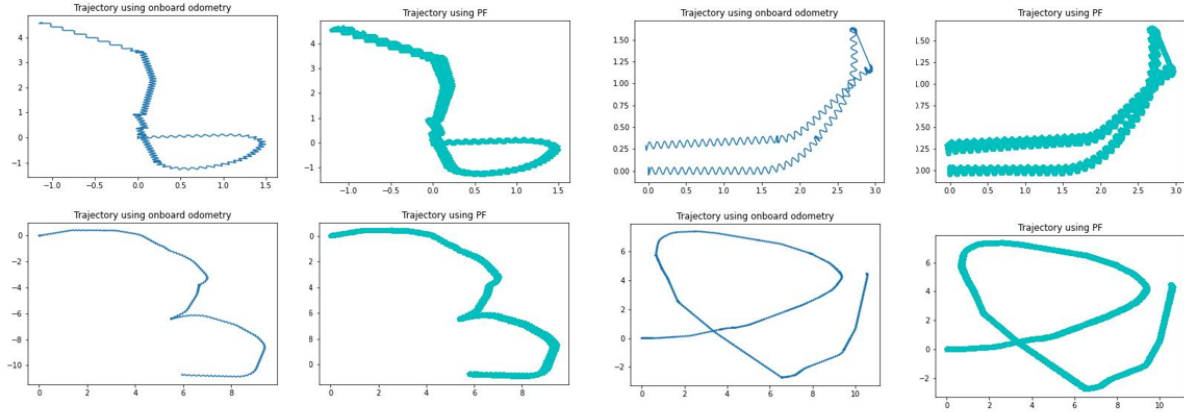


Figure 3: Odometry trajectory particle trajectories

2.2 (f) The full SLAM algorithm




Figure 4: The final binarized version of the map

3 Problem 3

3.1 (a) Data Loading and COLMAP




Figure 5: Data Loading and COLMAP

$$H = 100 \tag{1}$$

$$w = 100 \tag{2}$$

$$w = 95.8 \tag{3}$$

```
1 def load_colmap_data():
2     r"""
3     After using colmap2nerf.py to convert the colmap intrinsics and extrinsics,
4     read in the transform_colmap.json file
5
6     Expected Returns:
7         An array of resized imgs, normalized to [0, 1]
8         An array of poses, essentially the transform matrix
9         Camera parameters: H, W, focal length
10
11     NOTES:
12         We recommend you resize the original images from 800x800 to lower resolution,
13         i.e. 200x200 so it's easier for training. Change camera parameters accordingly
14     """
15     ##### YOUR CODE START #####
16
17     json_path='transforms_colmap.json'
18     image_dir='./data/data/images'
19     resize_dim=(100, 100)
20
```

```

21 # Load JSON file
22 with open(json_path, 'r') as f:
23     data = json.load(f)
24
25 # Initialize lists to store processed data
26 imgs = []
27 poses = []
28
29 # Process each frame in the JSON file
30 for frame in data['frames']:
31     # Load and resize image
32     img_path = os.path.join(image_dir, frame['file_path'][0])
33     img = Image.open(img_path).resize(resize_dim)
34     imgs.append(np.array(img) / 255.0) # Normalize to [0, 1]
35
36     # Process pose
37     pose = np.array(frame['transform_matrix'])
38     #print (pose)
39     poses.append(pose)
40
41 # Assuming all images have the same size and focal length
42 H, W = resize_dim
43 focal = frame["camera_angle_x"] * W / (2 * np.tan(frame['camera_angle_x'] / 2))
44
45 # Convert lists to numpy arrays
46 imgs = np.array(imgs)
47 poses = np.array(poses)
48
49 # Camera parameters: Height, Width, Focal length
50 camera_params = (H, W, focal)
51
52 return imgs, poses, camera_params

```

Code Listing 4: Calculate Camera Parameters

3.2 (b) Implementation of the NeRF

```
1 # -*- coding: utf-8 -*-
2 """Zhanqian Wu ESE 650 HW3 P3 Nerf
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/19y_9R6bB_r93PXfga6C_vNJZMRBfSzBU
8 """
9
10 from typing import Optional, Tuple
11 import torch
12 from torch import nn
13 from torch.nn import functional as F
14 import numpy as np
15 import matplotlib.pyplot as plt
16 from tqdm import tqdm
17 import os
18 import json
19 #import imageio
20 import cv2
21 import os
22 os.environ['PYTORCH_CUDA_ALLOC_CONF'] = 'expandable_segments:True'
23 from PIL import Image
24
25 from google.colab import drive
26 drive.mount('/content/drive')
27
28 def load_colmap_data():
29     r"""
30     After using colmap2nerf.py to convert the colmap intrinsics and extrinsics,
31     read in the transform_colmap.json file
32
33     Expected Returns:
34         An array of resized imgs, normalized to [0, 1]
35         An array of poses, essentially the transform matrix
36         Camera parameters: H, W, focal length
37
38     NOTES:
39         We recommend you resize the original images from 800x800 to lower resolution,
40         i.e. 200x200 so it's easier for training. Change camera parameters accordingly
41     """
42     ##### YOUR CODE START #####
43
44     json_path='transforms_colmap.json'
45     image_dir='./data/data/images'
46
47     json_path='/content/drive/MyDrive/ESE 650_HW3_P3/transforms_colmap.json'
48     image_dir='/content/drive/MyDrive/ESE 650_HW3_P3/data/data/images'
49     resize_dim=(200, 200)
50
51     # Load JSON file
52     with open(json_path, 'r') as f:
53         data = json.load(f)
54
55     # Initialize lists to store processed data
56     imgs = []
57     poses = []
58
59     # Process each frame in the JSON file
60     for frame in data['frames']:
61         # Load and resize image
62         img_path = os.path.join(image_dir, frame['file_path'][0])
```

```

63     img = Image.open(img_path).resize(resize_dim)
64     imgs.append(np.array(img) / 255.0) # Normalize to [0, 1]
65
66     # Process pose
67     pose = np.array(frame['transform_matrix'])
68     #print (pose)
69     poses.append(pose)
70
71     # Assuming all images have the same size and focal length
72     H, W = resize_dim
73     focal = frame["camera_angle_x"] * W / (2 * np.tan(frame['camera_angle_x'] / 2))
74
75     # Convert lists to numpy arrays
76     imgs = np.array(imgs)
77     poses = np.array(poses)
78
79     # Camera parameters: Height, Width, Focal length
80     camera_params = (H, W, focal)
81
82     return imgs, poses, camera_params
83     ##### YOUR CODE END #####
84
85
86 def get_rays(H, W, focal, c2w,):
87     """
88     Compute rays passing through each pixel in the world frame.
89
90     Parameters:
91     - H: Height of the image.
92     - W: Width of the image.
93     - focal: Camera intrinsic matrix of shape (3, 3).
94     - c2w: Camera-to-world transformation matrix of shape (4, 4).
95
96     Returns:
97     - ray_origins: Array of shape (H, W, 3) denoting the origins of each ray.
98     - ray_directions: Array of shape (H, W, 3) denoting the direction of each ray.
99     """
100     # Generate mesh grid for pixel coordinates
101
102     # Detect if a GPU is available and choose the device accordingly
103     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
104     #device = 'cpu' #Due to OOM
105     #print(device)
106
107     # Ensure c2w is a torch.Tensor and move it to the chosen device
108     if not isinstance(c2w, torch.Tensor):
109         c2w = torch.from_numpy(c2w).float() # Convert from numpy to tensor if necessary
110     c2w = c2w.to(device)
111
112     # Generate a grid of (i, j) coordinates
113     i, j = torch.meshgrid(torch.linspace(0, W-1, W, device=device), torch.linspace(0, H-1, H,
114     device=device))
115     i = i.t().flatten()
116     j = j.t().flatten()
117
118     # Normalize pixel coordinates (assuming the image center as origin)
119     dirs = torch.stack([(i - W * 0.5) / focal, -(j - H * 0.5) / focal, -torch.ones_like(i, device=
120     device)], -1)
121
122     # Rotate ray directions from camera frame to the world frame
123     rays_d = torch.sum(dirs[... , None, :] * c2w[:3, :3], axis=-1)
124
125     # The origin of all rays is the camera position in the world frame
126     rays_o = c2w[:3, -1].expand(rays_d.shape)

```

```

126 # Reshape rays_o and rays_d to [H, W, 3]
127 rays_o = rays_o.view(H, W, 3)
128 rays_d = rays_d.view(H, W, 3)
129
130 return rays_o, rays_d
131
132 def sample_points_from_rays(ray_origins, ray_directions, snear, sfar, Nsample):
133     """
134     Sample 3D points along rays within the specified near and far bounds.
135
136     Parameters:
137     - ray_origins: Array of shape (H, W, 3) denoting the origins of each ray.
138     - ray_directions: Array of shape (H, W, 3) denoting the direction of each ray.
139     - snear: Scalar or array defining the near clipping distance for each ray.
140     - sfar: Scalar or array defining the far clipping distance for each ray.
141     - Nsample: Number of points to sample along each ray.
142
143     Returns:
144     - sampled_points: Array of shape (H, W, Nsample, 3) with sampled 3D points.
145     - depth_values: Array of shape (H, W, Nsample) with depth values of sampled points.
146     """
147
148     # Make sure snear and sfar are tensors
149     H, W, _ = ray_origins.shape
150     device = ray_origins.device
151
152
153     # Compute the depth values for each sample
154     depth_values = torch.linspace(snear, sfar, Nsample, device=device)
155
156     depth_values = depth_values.expand(H, W, Nsample) # Make sure depth values have shape (H, W,
157     Nsample)
158
159     # Compute the 3D positions of each sample point along the rays
160     sampled_points = ray_origins[..., None, :] + depth_values[..., :, None] * ray_directions[...,
161     None, :] # Correct shape: (H, W, Nsample, 3)
162
163     return sampled_points, depth_values
164
165 def positional_encoding(x, max_freq_log2=10, include_input=True):
166     """Apply positional encoding to the input. (Section 5.1 of original paper)
167     We use positional encoding to map continuous input coordinates into a
168     higher dimensional space to enable our MLP to more easily approximate a
169     higher frequency function.
170
171     Expected Returns:
172     pos_out: positional encoding of the input tensor.
173             (H*W*num_samples, (include_input + 2*freq) * 3)
174     """
175     frequencies = 2 ** torch.linspace(0, max_freq_log2, steps=max_freq_log2+1, device=x.device)
176     # Create a list of frequencies, (sin(2^k * x), cos(2^k * x)) for k=0,...,max_freq_log2
177     encodings = [torch.sin(x * freq) for freq in frequencies] + [torch.cos(x * freq) for freq in
178     frequencies]
179     # Stack all encodings along the last dimension
180     encoded = torch.cat(encodings, dim=-1)
181
182     if include_input:
183         # Concatenate the original input with the encoded features
184         pos_out = torch.cat([x, encoded], dim=-1)
185     else:
186         pos_out = encoded
187
188     return pos_out

```



```

188
189
190 def volume_rendering(
191     radiance_field: torch.Tensor,
192     ray_origins: torch.Tensor,
193     depth_values: torch.Tensor
194 ) -> Tuple[torch.Tensor]:
195     """
196     Differentiably renders a radiance field, given the origin of each ray in the bundle,
197     and the sampled depth values along them.
198
199     Args:
200         radiance_field: Tensor containing RGB color and volume density at each query location,
201                         shape (H, W, num_samples, 4).
202         ray_origins: Origin of each ray, shape (H, W, 3).
203         depth_values: Sampled depth values along each ray, shape (H, W, num_samples).
204
205     Returns:
206         rgb_map: Rendered RGB image, shape (H, W, 3).
207     """
208     # Extract sigma (density) and color from the radiance field
209     sigma = torch.relu(radiance_field[..., 3]) # Extract volume density
210     rgb = torch.sigmoid(radiance_field[..., :3]) # Extract RGB colors
211
212     # Compute depth intervals
213     dists = torch.cat([depth_values[..., 1:] - depth_values[..., :-1], torch.tensor([1e10], device=
214 device).expand(depth_values[..., :1].shape)], dim=-1)
215     alpha = 1.0 - torch.exp(-sigma * dists)
216     weights = alpha * torch.cumprod(torch.cat([torch.ones_like(alpha[..., :1]), 1.0 - alpha + 1e
217 -10], dim=-1), dim=-1)[..., :-1]
218
219     rgb_map = torch.sum(weights[..., None] * rgb, dim=-2)
220
221     return rgb_map
222
223 class TinyNeRF(torch.nn.Module):
224     def __init__(self, pos_dim, fc_dim=128):
225         r"""Initialize a tiny nerf network, which composed of linear layers and
226         ReLU activation. More specifically: linear - relu - linear - relu - linear
227         - relu -linear. The module is intentionally made small so that we could
228         achieve reasonable training time
229
230         Args:
231             pos_dim: dimension of the positional encoding output
232             fc_dim: dimension of the fully connected layer
233         """
234         super().__init__()
235
236         self.nerf = nn.Sequential(
237             nn.Linear(pos_dim, fc_dim),
238             nn.ReLU(),
239             nn.Linear(fc_dim, fc_dim),
240             nn.ReLU(),
241             nn.Linear(fc_dim, fc_dim),
242             nn.ReLU(),
243             nn.Linear(fc_dim, 4)
244         )
245
246     def forward(self, x):
247         r"""Output volume density and RGB color (4 dimensions), given a set of
248         positional encoded points sampled from the rays
249         """
250         x = self.nerf(x)
251         return x

```

```

251
252 def get_minibatches(inputs: torch.Tensor, chunksize: Optional[int] = 1024 * 8):
253     r"""Takes a huge tensor (ray "bundle") and splits it into a list of minibatches.
254     Each element of the list (except possibly the last) has dimension 0 of length
255     chunksize.
256     """
257     return [inputs[i:i + chunksize] for i in range(0, inputs.shape[0], chunksize)]
258
259
260 def nerf_step_forward(height, width, focal_length, trans_matrix,
261                       near_point, far_point, num_depth_samples_per_ray,
262                       get_minibatches_function, model):
263     r"""Perform one iteration of training, which take information of one of the
264     training images, and try to predict its rgb values
265
266     Args:
267         height: height of the image
268         width: width of the image
269         focal_length: focal length of the camera
270         trans_matrix: transformation matrix, which is also the camera pose
271         near_point: threshold of nearest point
272         far_point: threshold of farthest point
273         num_depth_samples_per_ray: number of sampled depth from each rays in the ray bundle
274         get_minibatches_function: function to cut the ray bundles into several chunks
275         to avoid out-of-memory issue
276
277     Expected Returns:
278         rgb_predicted: predicted rgb values of the training image
279     """
280     ##### YOUR CODE START #####
281
282     # Step 1: Generate rays
283     # an implementation of get rays function that returns ray origins and directions
284     ray_origins, ray_directions = get_rays(height, width, focal_length, trans_matrix)
285
286
287     # Step 2: Sample points along each ray
288     # an implementation of sample points from rays that returns sampled points and depth values
289     sampled_points, depth_values = sample_points_from_rays(ray_origins, ray_directions, near_point,
290                                                            far_point, num_depth_samples_per_ray)
291
292
293     # Step 3: Apply positional encoding
294     # positional encoding expects a flattened list of points
295     flattened_sampled_points = sampled_points.reshape(-1, 3) # Flattening sampled points for
296     # positional encoding
297     positional_encoded_points = positional_encoding(flattened_sampled_points) # Apply positional
298     # encoding
299     # print("positional_encoded_points",positional_encoded_points.shape)
300
301
302     # Step 4: Run the model in batches
303     # Splitting the points into manageable chunks to avoid OOM
304     batches = get_minibatches_function(positional_encoded_points, chunksize=16384)
305     predictions = []
306     for batch in batches:
307         #print(batch.shape)
308         predictions.append(model(batch))
309
310     radiance_field_flattened = torch.cat(predictions, dim=0)
311
312     ##### YOUR CODE START #####
313
314     # Step 5: Volume rendering
315     # Reshape the radiance field to its unflattened shape

```

```

313 radiance_field = radiance_field_flattened.view(height, width, num_depth_samples_per_ray, 4)
314 #print(f"radiance_field is stored on: {radiance_field.device}")
315
316
317 # volume rendering that takes the radiance field, ray origins, and depth values
318 rgb_predicted = volume_rendering(radiance_field, ray_origins, depth_values)
319 #print("rgb_predicted",rgb_predicted.shape)
320
321
322 return rgb_predicted
323
324
325
326 def train(images, poses, hwf, near_point,
327          far_point, num_depth_samples_per_ray,
328          num_iters, model, History, DEVICE="cuda"):
329     r"""Training a tiny nerf model
330
331     Args:
332         images: all the images extracted from dataset (including train, val, test)
333         poses: poses of the camera, which are used as transformation matrix
334         hwf: [height, width, focal_length]
335         near_point: threshold of nearest point
336         far_point: threshold of farthest point
337         num_depth_samples_per_ray: number of sampled depth from each rays in the ray bundle
338         num_iters: number of training iterations
339         model: predefined tiny NeRF model
340     """
341
342
343     H, W, focal_length = hwf
344     H = int(H)
345     W = int(W)
346     n_train = images.shape[0]
347
348     # Optimizer parameters
349     lr = 5e-3
350     optimizer = torch.optim.Adam(model.parameters(), lr=lr)
351
352     # Seed RNG, for repeatability
353     seed = 9458
354     torch.manual_seed(seed)
355     np.random.seed(seed)
356
357
358     for _ in tqdm(range(num_iters)):
359         # Randomly pick a training image as the target, get rgb value and camera pose
360         train_idx = np.random.randint(n_train)
361         train_img_rgb = images[train_idx, ..., :3]
362         train_pose = poses[train_idx]
363
364         # Run one iteration of TinyNeRF and get the rendered RGB image.
365         rgb_predicted = nerf_step_forward(H, W, focal_length,
366                                         train_pose, near_point,
367                                         far_point, num_depth_samples_per_ray,
368                                         get_minibatches, model)
369
370
371         train_img_rgb_tensor = torch.from_numpy(train_img_rgb)
372         train_img_rgb = train_img_rgb_tensor.to(DEVICE).to(dtype=torch.float32)
373
374
375         # Compute mean-squared error between the predicted and target images
376         loss = torch.nn.functional.mse_loss(rgb_predicted, train_img_rgb)
377         loss.backward()

```

```

378     optimizer.step()
379     optimizer.zero_grad()
380     History.append(loss.cpu().detach().numpy())
381
382
383     if _ % 100 == 0:
384         with torch.no_grad():
385             plt.figure(figsize=(10, 4))
386
387             plt.subplot(121)
388             plt.imshow(train_img_rgb_tensor.cpu().numpy())
389             plt.title(f"Target Image at Iteration {_}")
390
391             plt.subplot(122)
392             plt.imshow(rgb_predicted.detach().cpu().numpy())
393             plt.title(f"Predicted Image at Iteration {_}")
394
395             plt.show()
396
397             print('Loss at '+str(_)+" iterations: ",loss.cpu().detach().numpy())
398
399     print('Finish training')
400
401 if __name__ == "__main__":
402     images, poses, hwf=load_colmap_data()
403     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
404     #device = 'cpu' #Due to OOM
405     near_point=2.
406     far_point=6.
407     num_depth_samples_per_ray = 64
408     num_iters = 2000
409     model = TinyNeRF(69)
410     model.to(device)
411
412     History=[]
413     train(images, poses, hwf, near_point,
414           far_point, num_depth_samples_per_ray,
415           num_iters, model, History,DEVICE=device)
416
417 plt.plot(History)
418 plt.grid()
419 plt.xlabel("Epoch")
420 plt.ylabel("Loss")
421 plt.title("ESE 650 Zhanqian Nerf Training Process")
422 plt.show()

```

Code Listing 5: NERF Code

3.3 (c) Network Training

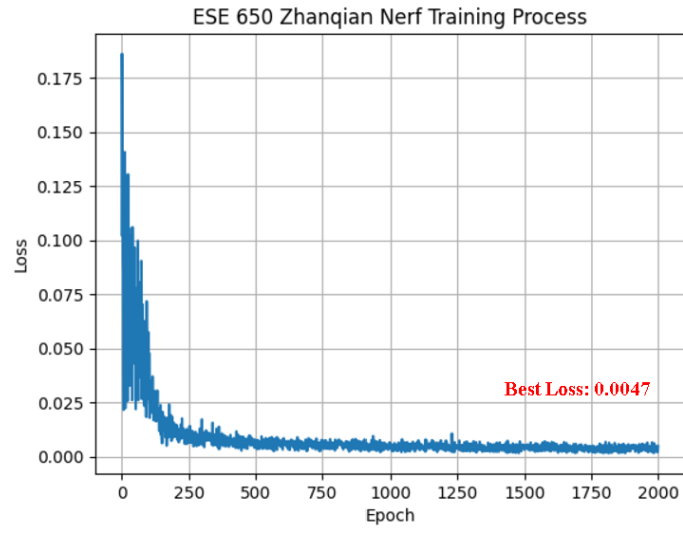


Figure 6: Nerf training loss graph

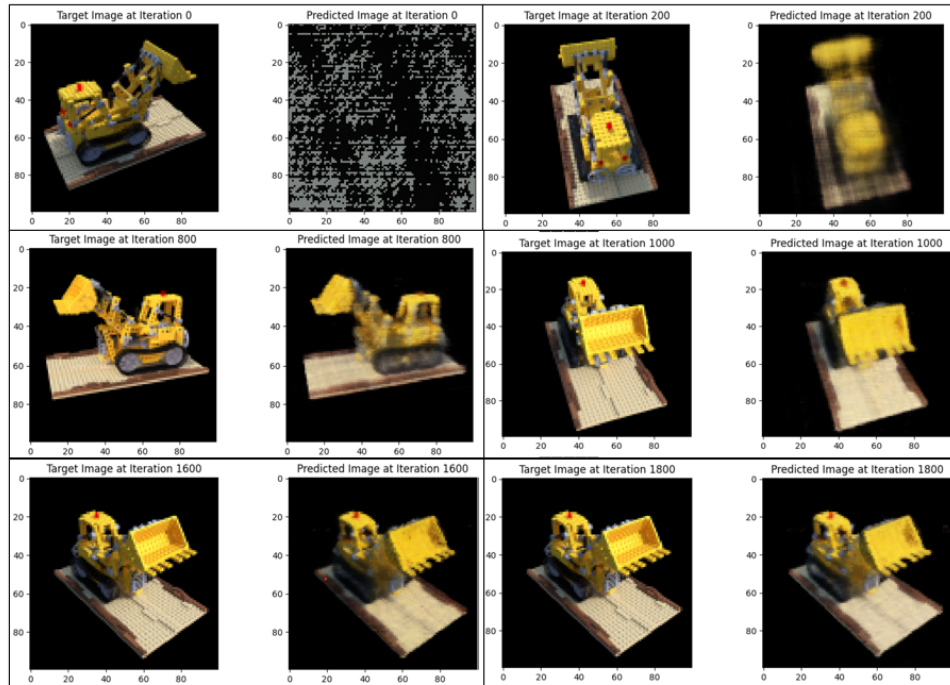


Figure 7: Comparison of the Results with the Number of Training Iterations

3.4 (d) Inference




Figure 8: Five viewpoints and rendered RGB images

Where NeRF is Working Well:

- **Overall Shape Recognition:** In several predicted images, NeRF seems to capture the overall shape of the object quite well. The contours and the bulk of the object are recognizable.
- **Color Approximation:** The colors in the predicted images, while sometimes a bit off, are relatively close to the target images, indicating that NeRF is effectively capturing the color information.

Where NeRF is Struggling:

- **Fine Details:** NeRF appears to struggle with fine details. In the predicted images, areas with small or intricate details are blurred or misrepresented, such as the precise edges and small features of the object.
- **Texture Fidelity:** The textures in the predicted images are less crisp and detailed compared to the target images. This can be observed in the loss of texture clarity on surfaces.
- **Edges and Boundaries:** The edges of the objects in the predicted images are not as sharp as in the target images, showing some challenges in rendering sharp boundaries.
- **Consistent Lighting and Shadows:** Some of the predicted images appear to have slightly different lighting conditions and shadowing compared to the target images, suggesting difficulty in capturing the exact lighting of the scene.

In summary, Neural Radiance Fields often struggle with high-frequency details due to several factors inherent in their design and the training data used. High-frequency details refer to rapid changes in color or brightness in an image, which correspond to the fine details in a scene, like textures or edges.

References

- [1] hmmmlearn/hmmmlearn. (2022, November 25). GitHub. <https://github.com/hmmmlearn/hmmmlearn>